

FAYAZ

Lead Full-Stack TypeScript Engineer + AI / Retrieval Engineer

Proposed for Roles 1 & 3 — Mubarak AI Minimum Retained Engineering Squad

mfaiyazuddeen@gmail.com
Chennai, India (GST, UTC+5:30)
10+ yrs production engineering · AWS Certified
Solutions Architect
portfolio site on file

One engineer, two roles, zero overlap in hours billed: a decade of production full-stack ownership paired with hands-on RAG systems already shipping grounded, multilingual AI tutoring in production — the closest possible precedent to Mubarak AI's own tutoring platform.

TRACK RECORD, IN NUMBERS

10+ years production engineering	100+ developers mentored	8+ international client engagements	71% form load time cut (14s → 4s)	90% data upload time cut (300min → 30min)
120% render throughput gain, zero legacy code touched	100% offline coverage delivered, zero API code changed	85% UI interaction latency cut (6s → <1s)	0 browser memory crashes after fix (10yr datasets)	6 languages supported in a live AI tutoring platform

FIT FOR ROLE 1 — LEAD FULL-STACK TYPESCRIPT ENGINEER

10+ years owning full-stack delivery end-to-end — UI, server APIs, integrations, and production incident response — across 8+ international client engagements (US, Canada, Denmark, Australia).

- Led technical delivery as Technical Lead across a 5–10 engineer team: architecture decisions, code review, mentoring, release ownership.
- Migrated 50+ legacy data grids to modern React on a US compliance platform (Averq) while the product stayed live — no downtime, no rewrite.
- Shipped real-time collaboration and offline-first reliability: WebSocket conflict detection plus a PWA offline layer covering 100% of sessions with zero lines of existing API code changed.
- Built reactive Kotlin/Spring Boot services and a high-performance GTFS cache for a national transit platform (Nexum) handling high-volume, high-availability traffic.
- Rebuilt a 3,500-question form engine's evaluation pipeline, cutting load time 71% (14s → 4s) — now the platform's standard pattern for all new forms.
- Comfortable owning day-to-day technical leadership as the single point of contact: turning requirements into tasks, coordinating releases across specialists, staying on defects through verified resolution.

FIT FOR ROLE 3 — AI & RETRIEVAL ENGINEER

Currently the AI Product Engineer building **Robogebra**, a production AI math tutoring platform — the same problem shape as Mubarak AI: grounded student-facing answers, not a generic chatbot.

- Built the full RAG pipeline in production with LangChain + ChromaDB, grounding every student answer in trusted study material instead of open-domain model knowledge.
- Built hybrid retrieval (vector + keyword) in MongoDB Atlas so students get relevant explanations even from vague or partial questions — directly applicable to pgvector tuning at Mubarak AI's scale.
- Shipped a document ingestion pipeline for PDFs — upload to searchable, chunked, embedded content — the same OCR-to-embeddings path Mubarak AI's curriculum ingestion needs.
- Delivered real-time streaming answers via Server-Sent Events, and a safe KaTeX rendering pipeline for math notation, so answers appear instantly without broken output in production.

- Designed multilingual response flows across 6 language modes — direct precedent for Mubarak AI's Arabic/English grounded-tutoring requirement.
- Production-first mindset on AI safety and cost: evaluating groundedness, output safety, and provider spend are daily concerns, not one-off exercises.

WHY ONE PERSON FOR BOTH ROLES

- No handoff loss: application architecture and retrieval/grounding decisions are made by the same person, so RAG changes never wait on a cross-team sync.
- Full allocation transparency: hours for Role 1 and Role 3 tracked and billed separately against the engagement, never double-counted.
- Backup and continuity covered by the wider delivery team already in place across 8+ prior client engagements — absence and handover is a practiced process, not a first attempt.

EXPERIENCE

AI Product Engineer — Provility Jan 2024 – Present

RAG pipelines, real-time LLM streaming (SSE), agentic tool-calling, multi-stage AI workflows in production.

Technical Lead — Provility Apr 2019 – Jan 2024

Led a cross-functional engineering team; owned architecture, code review, mentoring, delivery.

Senior Engineer — Provility Jun 2016 – Apr 2019

Designed backend systems and APIs; performance optimization; cross-functional delivery.

Junior Engineer — Provility Sep 2014 – Jun 2016

Web application features, QA-driven bug fixes, sprint delivery support.

CERTIFICATION

AWS Certified Solutions Architect

LOGISTICS (per vendor response requirements)

Working location	Chennai, India
Time zone	India Standard Time (UTC+5:30)
Support window coverage	Sun–Thu 07:00–22:00 AST overlaps IST afternoon/evening — exact shift and 24-hr P1 on-call pricing to be confirmed with vendor
Allocation to this engagement	To be confirmed — proposed across Role 1 (1.0 FTP) and Role 3 (1.0 FTP), hours tracked separately, no double-counting
Earliest start date	To be confirmed with vendor
Backup / continuity	Covered via delivery team; details to be finalized in vendor proposal